

# Relational Data Handling

Siddharth Kaza

Towson University, Computer and Information Sciences

# Outline

---

- ▶ Web Applications and Relational Data
  - ▶ Introduction
  - ▶ Simple CRUD operations in ActiveRecord (Rail's ORM)
- ▶ Relationships between models/tables
- ▶ STI and Polymorphic associations
- ▶ Transactions

# Relational data

---

- ▶ Relational data – model relations between items
  - ▶ Containment – one:many
    - ▶ cart contains many line\_items
  - ▶ Relationships – many:many
    - ▶ Authors have many books
    - ▶ Books have many authors
- ▶ How do we do this?

# Traditional approach

---

1. Model things in SQL
2. Every table has a primary key
  1. Identifier that uniquely names each row
  2. Try to maintain normal forms
  3. Include foreign keys in a table to indicate a relationship
3. Use SQL queries to create associations
4. Write code to process rows in query result.

# 1. Model in SQL

---

- ▶ Consider a book store

- ▶ Table definitions with primary keys

```
create table books (book_id int, title varchar(80), publisher_id int)
```

```
create table publishers (publisher_id int, name varchar(80))
```

```
create table authors (author_id,int, first_name varchar(20),  
                    last_name varchar(30))
```

```
create table authors_books (author_id int, book_id int)
```

Primary keys



Foreign keys



- ▶ Normal forms

- ▶ Each column is atomic

- ▶ Other columns must provide facts about the primary key and nothing but the primary key

## 2. SQL Query and 4. processing result

---

- ▶ query = “Select b.title, p.name, a.first\_name, a.last\_name from book b, publishers p, authors a, authors\_books ab where b.publisher\_id=p.publisher\_id and ab.book\_id=b.book\_id and ab.authors\_id = a.authors\_id and b.book\_id=36”;
- ▶ ResultSet res = stmt.executeQuery(query)  
while (res.next) {  
    String booktitle = rs.getString(1);  
    String pubname = rs.getString(2);  
    ...  
}
- ▶ stuff things into variables, classes, objects, etc.

# What's the problem with this?

---

- ▶ SQL is painful to write
  - ▶ inside java/C#/... code, DB enthusiasts like to write stored procedures outside
- ▶ Follow-on code must be modified if query changes
- ▶ Am I pulling in too much data? Too little?
- ▶ Relationships are not easily discernible from code
- ▶ Details in too many places
- ▶ => Need for Object-Relational Mapping

# CRUD

---

- ▶ ORM should be able to provide the basic DB operations
- ▶ Create
- ▶ Read
- ▶ Update
- ▶ Delete
- ▶ Let look at these from the depot context
- ▶ Find good details here:  
[http://guides.rubyonrails.org/active\\_record\\_querying.html](http://guides.rubyonrails.org/active_record_querying.html)



# Create

---

▶ New creates new object

▶ Save saves it

```
an_order = Order.new
```

```
an_order.name = "Dave Thomas"
```

```
an_order.email = "dave@pragprog.com"
```

```
an_order.address = "123 Main St"
```

```
an_order.pay_type = "check"
```

```
an_order.save
```

# More create

---

## ► Use hash (useful to get data from forms)

```
an_order = Order.new(  
  :name => "Dave Thomas" ,  
  :email => "dave@pragprog.com" ,  
  :address => "123 Main St" ,  
  :pay_type => "check" )  
an_order.save
```

## ► use block

```
user = User.new do |u|  
  u.name = "David"  
  u.occupation = "Code Artist"  
end
```

# save vs. save!, create vs. create!

---

- ▶ `save` returns true if record saved, nil otherwise
  - ▶ Nil if validation errors
  - ▶ Assume coming from a form that will show errors
- ▶ `create` returns object whether saved or not
  - ▶ You must check the validations
- ▶ `save!`, `create!` Raise exception
  - ▶ Use this to force you to write code that handles errors

if order.save

# all OK

else

# validation failed

end

# READ - Find

---

`p = Product.find(1)`

- ▶ Every model class supports the find method
- ▶ Takes one or more primary keys
- ▶ If given one
  - ▶ Returns object for that row
- ▶ If multiple – returns array
- ▶ RecordNotFound if any not found
- ▶ `Product.find(27)`
  - ▶ `Select * from product where product_id=27`

`Product.find([27,34,59])`

- ▶ `Select * from product where product_id in (27,34,59);`

# Find all, first

---

- ▶ **Product.all**
  - ▶ Finds all
- ▶ **Product.first**
  - ▶ Finds first
  - ▶ Same query, but *first* just returns one row
- ▶ **Subject to conditions**
  - ▶ `pos = Order.where("name = 'Dave' and pay_type = 'po'")`
- ▶ **Essentially generates arbitrary conditions in SQL**
  - ▶ The string after condition is just appended to the generated statement

# Warning on using :conditions

---

- ▶ Parameterize them – not a hard-wired query

```
pos = Order.where("name = '#{name}' and pay_type =  
                'po'")
```

- ▶ Not a good idea! – if name is entered from a form
- ▶ SQL injection – can open up security holes
  - ▶ consider this: the user enters “john OR 1--” in the name field
  - ▶ “OR 1” will make everything true
  - ▶ “--” will comment everything after that

# Warning on using :conditions (cont.)

---

- ▶ Use ? as a placeholder

```
pos = Order.where ("name = ? and pay_type = 'po'" ,  
name)
```

- ▶ Or hash (with named parameters)

```
pos = Order.where("name = :name and pay_type = :pay_type"  
  ,{:pay_type => pay_type, :name => name})
```

- ▶ When you use parameters - '?', Active Record will quote all characters that have special meaning for that DB adapter.

# Warning on using :conditions (cont.)

---

- ▶ If you need to execute query with similar options several places in your code
  - ▶ then put it as a method in the model

Instead of

```
Email.where("user_id = ? and read='no'",  
            user.id)
```

Do this

```
class Email < ActiveRecordBase  
  def self.find_unread_email_for_user (user)  
    Email.where("user_id = ? and  
                read='no'", user.id)  
  end  
end
```



# Send params for query as a hash

---

- ▶ Suppose form asks for name and pay type in an Order object

```
pos = Order.where("name = :name and pay_type = :pay_type" ,  
                  params[:order])
```

- ▶ Or

```
pos = Order.where(params[:order])
```

- ▶ This will search on all fields in form – not just name and pay\_type
  - ▶ It will use the 'and' operator between conditions in the where clause

# Other find options

---

- ▶ **:order**
- ▶ **:limit**
- ▶ **:offset**
  - ▶ ( for pagination, etc.)
- ▶ **:joins**
  - ▶ Model should cover this for you
- ▶ **:select**
  - ▶ Which columns – not just all
- ▶ **:readonly**
- ▶ **:group**
- ▶ **:lock**
  - ▶ Lock the row while you are working on it, DB specific

# Using native SQL in find

---

## ► find\_by\_sql

```
orders = LineItem.find_by_sql  
("select line_items.* "+  
" from line_items, orders " +  
" where order_id = orders.id " +  
" and orders.name = 'Dave Thomas' " )
```

# More finds – Aggregates

---

- ▶ `Order.average(:amount)`
- ▶ `Order.maximum(:amount)`
- ▶ `Order.minimum(:amount)`
- ▶ `Order.sum(:amount)`
  
- ▶ `std_dev = Order.calculate(:std ,:amount)`
  - ▶ `std` as an SQL function
  - ▶ Options similar to others
  - ▶ Can be grouped

# Grouping and finds

---

▶ `result = Order.maximum :amount, :group => "state"`

`puts result #=> [{"TX", 12345}, {"NC", 3456}, ...]`

# Counts

---

- ▶ No parameters – all rows
  - ▶ `Order.count`
- ▶ With parameters
- ▶ `result = Order.count "amount > 10"`
- ▶ By sql

```
count = LineItem.count_by_sql(  
  "select count(*) " +  
  "from line_items, orders " +  
  "where line_items.order_id = orders.id " +  
  "and orders.name = 'Dave Thomas' " )
```

# Dynamic finders

---

- ▶ Order – attributes name, email, address

```
order = Order.find_by_name("Dave Thomas" )
orders = Order.find_all_by_name("Dave Thomas" )
order = Order.find_all_by_email(params['email' ])
```
- ▶ find\_by\_xxx is :first
- ▶ find\_all\_by\_xxx is :all
- ▶ Conjunctions (deprecated in Rails 4, gem available)

```
user = User.find_by_name_and_password(name, pw)
```
- ▶ Combine create and find – create if not there

```
cart = Cart.find_or_initialize_by_user_id(user.id)
cart.save
```

# Update

---

- ▶ Simply modify and save

```
order = Order.find(1)
order.quantity = order.quantity + 1
order.save
```

- ▶ `update_attributes`

```
order = Order.find(321)
order.update_attributes(:name => "Barney" , :email =>
"barney@bedrock.com" )
```

- ▶ Combine with `find`

```
▶ order = Order.update(12, :name => "Barney" , :email =>
"barney@bedrock.com" )
```



# Delete

---

## `order.destroy`

- ▶ Deletes row from db
- ▶ Freezes object from future changes

## `order.destroy_all` ([“shipped\_at < ?”, 30.days.ago])

- ▶ Destroy all order object
- ▶ The `destroy` method is a ORM level method
  - ▶ It takes into consideration all validations, relationships, and callbacks
- ▶ There is another method called `delete` (`order.delete`)
  - ▶ This is a pure SQL method and not recommended

# Relationships

---

- ▶ How do we define them? – foreign keys
- ▶ Pros:
  - ▶ Forces database to do some work for you
  - ▶ Gives us explicit relationships
- ▶ Cons:
  - ▶ RDMBS-specific
  - ▶ May be hard to debug when things go wrong
- ▶ Add them (optionally) upon deployment?

# Defining Relationships in ORM

---

- ▶ There is the DB level
  - ▶ foreign keys

- ▶ and then there is the ORM/Rails level.

Using the Depot example from the book:

```
class Order < ActiveRecord::Base
  has_many :line_items
```

```
class Product < ActiveRecord::Base
  has_many :line_items
  # ...
end
```

# Relationships in Rails Models

---

```
class LineItem < ActiveRecord::Base
  belongs_to :order
  belongs_to :product
end
```

- ▶ tells rails that the rows in line\_items are children of rows in the orders and products table

- ▶ What do these do?

```
li = LineItem.find(...)
puts "This line item was bought by #{li.order.name} "
```

```
order = Order.find(...)
puts "This order has #{order.line_items.size} line items"
```

these are methods



# More on ActiveRecord and Relations

---

- ▶ Active record lets us naturally walk relationships in object-oriented manner

- ▶ Consider – each line\_item references a product

- ▶ No need to write SQL to join

```
select      p.price
from        line_items li,
            products p
where       li.product_id = p.id
            and p.id = 36
```

line_item_id	order_id	product_id
--------------	----------	------------

- ▶ Also, don't need to find second object

```
product_id = line_item.product_id
product = Product.find(product_id)
price = product.price
```

- ▶ use relations

```
price = line_item.product.price
```

# Primary Keys (aside)

---

- ▶ Primary key is the unique identifier for a database row
- ▶ Generally, but not always, an integer
- ▶ Rails really wants you to have an integer as a primary key
  - ▶ Rails wants to call it id
  - ▶ Usually a good idea
  - ▶ Even if you have some external id for something, having your own internal version might help
  - ▶ You can switch it off however -

```
create_table :categories_products, :id => false do |t|  
  t.column :product_id, :integer  
  t.column :category_id, :integer  
end
```

# Primary Keys (cont.)

---

- ▶ ISBNs, SSNs are supposedly unique
  - ▶ But may not be...
  - ▶ Also what if ISBN changes? – What would you have to do?
  - ▶ Why would you not use SSN?

```
class Book < ActiveRecord::Base
  self.primary_key = "isbn"
end
```

```
create_table :tickets, :primary_key => :number do |t|
  t.column :created_at, :timestamp
  t.column :description, :text
  t.column :number, :integer
end
```

# Naming tables

---

- ▶ Class names get split at caps and pluralized to form table names
  - ▶ Linetitem class becomes line\_items table
  - ▶ But, you use rails generate model line\_item
- ▶ Inherent tradeoff
  - ▶ Attempt at “smart” guessing of names simplifies coding, initial installation, etc.
- ▶ But, it may not always work the way you want
  - ▶ Might not always agree with the designer's decisions
- ▶ Alternative – configure it explicitly in files/code
  - ▶ Read through the configuration to figure it out
    - ▶ set\_table\_name “name” in model
    - ▶ or look at config/initializers/inflections.rb



# ActiveRecord – Relationships in detail

---

- ▶ **ActiveRecord supports three types of relationships**
  - ▶ one – many
  - ▶ many – many
  - ▶ one - one

# Setting up relationships

---

- ▶ Two pieces
  - ▶ Fields in table
  - ▶ Migration for model *yyy* has *xxx\_id* as a column if each *y* has a foreign key with *x*
    - ▶ ie., *x* is associated with many *y*'s
    - ▶ Or, *x* contains many *y*'s

```
class CreateLineItems < ActiveRecord::Migration
```

```
  def self.up
```

```
    create_table :line_items do |t|
```

```
      t.column :product_id, :integer, :null => false
```

```
      t.column :order_id, :integer, :null => false
```

```
      t.column :quantity, :integer, :null => false
```

```
      t.column :total_price, :decimal, :null => false
```

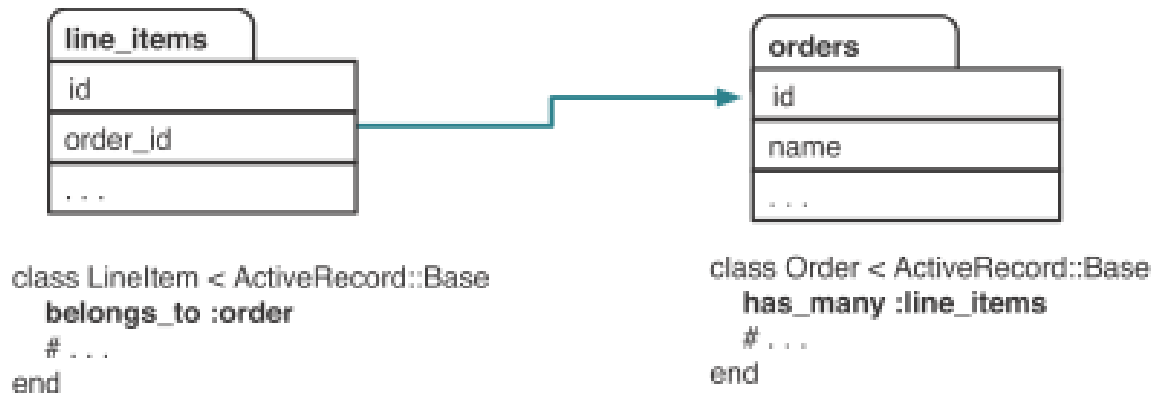
```
    end....
```

Named after  
the class not  
the table  
- singular!

# One – Many

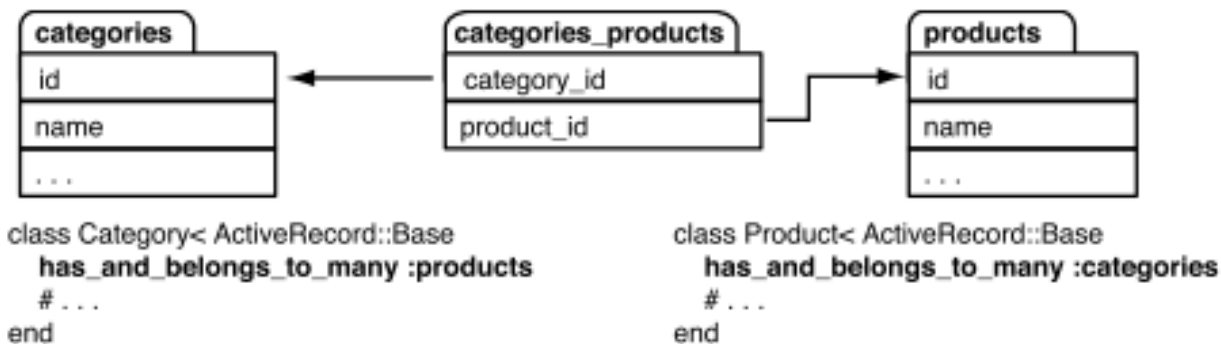
---

- ▶ Example – relationship between line\_items and orders



# Many – Many

- ▶ Example – Categories and Products – A product may have many categories and a category many have many products




- ▶ Within the db the many to many is represented by a join table – table1\_table2 (in alphabetical order)

# One - One

---

- ▶ Example – relationship between orders and invoices



- ▶ Model with the `foreign_key` has the `belongs_to` declaration

# Notes on this

---

- ▶ Foreign key named after class of table with `_id` appended
- ▶ Be careful of Rails' cleverness with plurals
  - ▶ Person model will map to people table, and `person_id` will be foreign key
- ▶ `has_many` takes full table name, as in database
- ▶ item with foreign key (`line_item`) always is the `belongs_to`
  - ▶ When you write *belongs\_to* think you are writing *references* in SQL terms

Declaration in child	Foreign Key	Parent Class	Parent Table
<code>belongs_to :product</code>	<code>product_id</code>	Product	products
<code>belongs_to :invoice_item</code>	<code>invoice_item_id</code>	InvoiceItem	invoice_items

# Methods that clauses provide us

---

- ▶ `belongs_to`

- ▶ `line_item belongs_to :product`
  - ▶ `line_item.product` – accessor and mutator
    - must save after mutate
- ▶ `build_product (attribute hash) coded as`
  - ▶ `line_item.product = Product.new (.....)`
    - create a new product
    - must save it

# Dependency options for has\_many

---

- ▶ `:dependent=>:destroy`
  - ▶ goes through child row and deletes each that refers to parent being deleted
- ▶ `:dependent=>:delete_all`
  - ▶ single sql query to delete all
- ▶ can override the search SQL by using the `:finder_sql` options
- ▶ can change the ordering
  - ▶ `:order => "quantity, unit_price DESC"`



# Methods you get with has\_many

---

- ▶ `order has_many line_items`
  - ▶ `line_items` – get array
  - ▶ `line_items << line_item`
  - ▶ `line_items.push(line_item)` – add to list
  - ▶ `line_items.replace(item1,item2...)`
    - ▶ replace old with new – detects differences and optimizes DB access
  - ▶ `line_items.delete(item1, items2...)`
    - ▶ invokes `destroy` as necessary
  - ▶ `line_items.delete_all`
  - ▶ `line_items.destroy_all`
  - ▶ `line_items.find`
  - ▶ `size, count, length` (forces reload), `empty`, `sum`, `uniq`, `build`, `create`

# has\_one options

---

- ▶ `:dependent=>:destroy` – delete the associated item in other table when this is destroyed
- ▶ `:dependent=>nullify` – child row is orphaned
- ▶ `:dependent=>false`
  - ▶ child is not changed.

# The :dependent option

```
class Order < ActiveRecord::Base
  has_one :invoice
end
```

orders		
id	name	...
300	Dave	
...	...	

invoices		
id	order_id	...
123	300	

order = Order.find(300)  
order.invoice = Invoice.new(...)

Existing invoice  
is orphaned

invoices		
id	order_id	...
123	NULL	
124	300	

Figure 19.1: Adding to a has\_one relationship

# A Real Estate DB Example – reflection on your lab

---

- ▶ 2 concepts
  - ▶ Home and owner
- ▶ Two tables
  - ▶ Homes
    - ▶ Street Number, Street, City, State, Zip
  - ▶ Owner
    - ▶ First, Last Name
- ▶ Foreign key
  - ▶ Owner has a home, or vice-versa

# Two-table design

---

- ▶ Suppose we want to have multiple owners for a given house
  - ▶ Owner has home\_id as a foreign key
  - ▶ Multiple owner records can refer to any given home
- ▶ Or, multiple homes for each owner
  - ▶ Home has owner\_id as foreign key
- ▶ What if we want both?

# Many-to-Many relationships: mapping tables

---

- ▶ A home can have many owners
- ▶ Owner can have many homes
- ▶ homes\_owners – mapping tables
  - ▶ 2 columns: owner\_id, home\_id
- ▶ One entry for each instance of a home owner
- ▶ Can add more attributes
  - ▶ Purchase date?
  - ▶ Primary residence? – boolean
  - ▶ Order of ownership?
  - ▶ How do you do this? More on this later....

# many-many relationships

---

- ▶ Home, owner related by has-many
  - ▶ homes\_owners join table

```
def self.up
```

```
  create_table :homes do |t|  
    t.column :street, :string
```

```
    ...
```

```
  end
```

```
  create_table :owners do |t|  
    t.column :firstname, :string
```

```
    ..
```

```
  end
```

```
  ...
```

# the join table

---

```
create_table homes_owners, :id => false do |t|  
  t.column :house_id, :integer  
  t.column :owner_id, :integer  
end
```

```
add_index :homes_owners, [:home_id, :owner_id]  
add_index :homes_owners, :owner_id
```

- ▶ join table is named after two tables it joins (names alphabetized)
- ▶ join table does not need a primary key
  - ▶ compound key from two columns
- ▶ indices help speed queries



# Indices

---

- ▶ Speed up database search
- ▶ Useful when we're dealing with something other than the primary key, and we're searching on it frequently
- ▶ In migration
  - ▶ `add_index :table, :column`
- ▶ For e.g., for the Orders model
  - ▶ `add_index :orders, :name`
- ▶ Generates SQL
  - ▶ `Create INDEX name_index on orders (name);`
- ▶ Composite indices – multiple columns – possible – pass column names as array
  - ▶ `add_index :names, [:lastname, :firstname]`

# specifying many-to-many in models

---

- ▶ `has_and_belongs_to_many`

```
class Home < ActiveRecord::Base
  has_and_belongs_to_many :owners
  ...
end
```

```
class Owner < ActiveRecord::Base
  has_and_belongs_to_many :homes
end
```

- ▶ Note that there is no model for `home_owners`

# has and belongs to many

---

- ▶ Owner has and belongs to many houses
- ▶ House has and belongs to many owners

Owner o = Owner.find(37)

houses = owner.houses

- ▶ Question – how do we handle attributes of this relationship?
  - ▶ purchase date
  - ▶ primary residence?
  - ▶ order – first, second, third, etc.

# Attributes of many-many relationships

---

- ▶ **Add columns to table**
  - ▶ one column for each new attribute
  - ▶ many-to-many table is just like a regular table
- ▶ **in Rails,**
  - ▶ we treat join tables separately
  - ▶ have a table declaration in the DB
  - ▶ but no ActiveRecord...
- ▶ **if we need attributes,**
  - ▶ define a new model – with active record – and use it

# many-many relationships with attributes

---

- ▶ Home and owner as before
  - ▶ define ownership table

```
create_table ownerships, :id => false do |t|  
  t.column :home_id, :integer  
  t.column :owner_id, :integer  
  t.column :purchase_date, :date  
  t.column :primary, :boolean  
  ...  
end
```

- ▶ Add attributes to table...

# Ownership ActiveRecord

---

```
class Ownership < ActiveRecord::Base  
  belongs_to :home  
  belongs_to :owner  
end
```

- ▶ both Home and Owner has\_many :ownerships

# using many-many relationships with attributes

---

```
ownership = Ownership.new  
ownership.purchase_date = Time.now  
ownership.primary = true  
ownership.owner = owner1  
ownership.house = house3  
ownership.save
```

- ▶ But, how can we get from owner to houses
  - ▶ and vice-versa?

# has\_many :through

---

- ▶ Used to indicate role of join classes that have active record implementations

```
class House < ActiveRecord::Base
  has_many :ownerships
  has_many :owners, :through => :ownerships
end
```

```
class Owner < ActiveRecord::Base
  has_many :ownerships
  has_many :homes, :through=> :ownerships
end
```

- ▶ now, we can say `home.owners << aNewOwner`
- ▶ etc.
- ▶ The `:through` option can also have `:conditions`



# Looking under the hood

---

- ▶ Efficient use of tools such as inferred relationships and accessors requires understanding of how the query generation works
- ▶ Put another way, this can lead to some woefully inefficient code if we're not careful
- ▶ Default in rails – defer loading child rows from the database until they are referenced
  - ▶ “lazy” loading

# Example of deferred loading

---

```
class Course < ActiveRecord::Base
  belongs_to :professor
  has_many :students
end
```

**Say you had to list all the course names, the professor teaching, and all the students**

```
for course in Course.find(:all)
  puts "Course: #{course.name}"
  puts "professor: #{course.professor.name}"
  for student in course.students
    puts "student: #{student.name}"
  end
end
```

- ▶ 1 query for find
- ▶ Then if there are  $n$  courses,  $n$  more for professor
- ▶  $n$  for students
- ▶ total of  $m$  students –  $m$  total for student name
- ▶  $2n+m+1$  queries
- ▶ So sometimes you might want to override deferred loading

# Alternative :include

---

```
for course in Course.find(:all, :include=> :professor)
  puts "Course: #{course.name}"
  puts "professor: #{professor.name}"
  for student in course.students
    puts "student: #{student.name}"
  end
end
```

- ▶ preloads professor
- ▶ pulls it in all at once
  - ▶  $n+m+1$  total

## Alternative :include (cont.)

---

```
for course in Course.find(:all, :include=>[ :professor,  
  :students])  
  puts "Course: #{course.name}"  
  puts "professor: #{professor.name}"  
  for student in course.students  
    puts "student: #{student.name}"  
  end  
end
```

► I query

# Caveats with :include

---

- ▶ May or may not improve performance
- ▶ May slow things down – lots of data
- ▶ Use it only if you really need all of that data
- ▶ Try both with and without – benchmark
- ▶ Be aware of what is being done when you chase too many relationships

# Look in the book/rails guides for finer points on

---

- ▶ **when items get saved**
  - ▶ save items manually and you won't hit problems
  - ▶ defensive programming
- ▶ **caching counters**
  - ▶ `order.line_items.size` will tell you how many
  - ▶ but there are ways to potentially speed this up
    - ▶ `:counter_cache => true` as a **`belongs_to`** option
    - ▶ `t.column :line_items_count, :integer, :default=>0` in the **orders migration**

# Inheritance and ORM

---

- ▶ Object-Oriented world makes great use of inheritance
- ▶ Example: roles of people
  - ▶ person: name, email
  - ▶ customer: person with a balance
  - ▶ employee: person with a boss and a department
  - ▶ manager: employee with a group and a list of managed employees,
  - ▶ etc...
- ▶ How do we use relational databases to model these inherited relationships?
- ▶ First – how would you make the relational schema to represent this?

## General idea: single-table inheritance (STI)

---

- ▶ all classes in hierarchy in one table
- ▶ column for each attribute of each class in hierarchy
- ▶ additional column - “type” - which type of object we're looking at



# STI- migration

---

```
create_table :people, :force => true do |t|
  t.column :type, :string

  # common attributes
  t.column :name, :string
  t.column :email, :string

  # attributes for type=Customer
  t.column :balance, :decimal, :precision => 10, :scale => 2

  # attributes for type=Employee
  t.column :reports_to, :integer
  t.column :dept, :integer

  # attributes for type=Manager
  # - none -
end
```

# Hierarchy of objects

---

```
class Person < ActiveRecord::Base  
end
```

```
class Customer < Person  
end
```

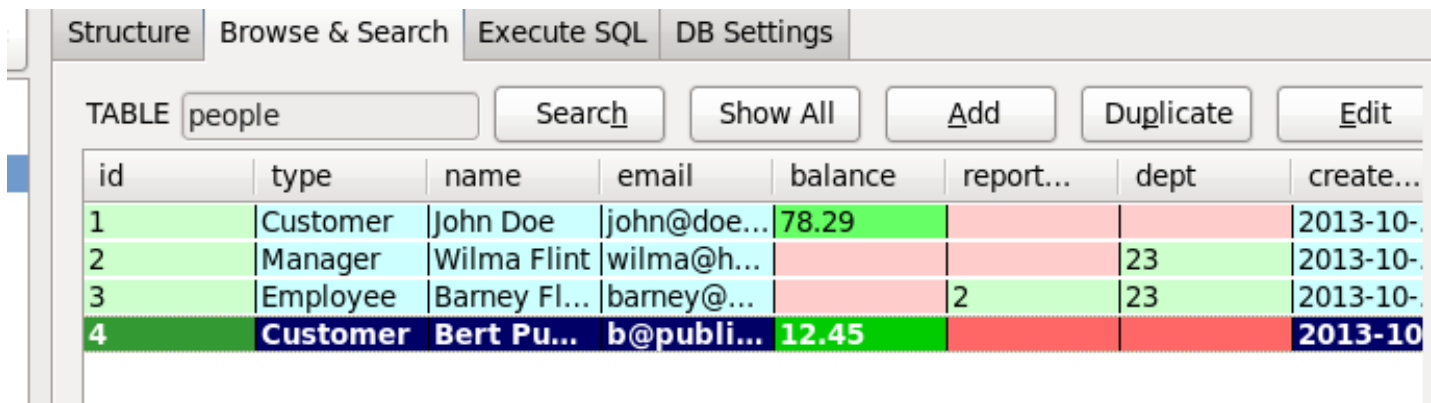
```
class Employee < Person  
  belongs_to :boss, :class_name => "Employee" , :foreign_key => :reports_to  
end
```

```
class Manager < Employee  
end
```

# STI in use

```
Customer.create(:name => 'John Doe' , :email => "john@doe.com" , :balance => 78.29)
wilma = Manager.create(:name => 'Wilma Flint' , :email => "wilma@here.com" , :dept => 23)
Customer.create(:name => 'Bert Public' , :email => "b@public.net" , :balance => 12.45)
barney = Employee.new(:name => 'Barney Rub' , :email => "barney@here.com" , :dept => 23)
```

```
barney.boss = wilma
barney.save!
```



id	type	name	email	balance	report...	dept	create...
1	Customer	John Doe	john@doe...	78.29			2013-10-
2	Manager	Wilma Flint	wilma@h...			23	2013-10-
3	Employee	Barney Fl...	barney@...		2	23	2013-10-
4	Customer	Bert Pu...	b@publi...	12.45			2013-10-

# STI in use

---

```
manager = Person.find_by_name("Wilma Flint" )  
puts manager.class #=> Manager  
puts manager.email #=> wilma@here.com  
puts manager.dept #=> 23
```

ActiveRecord gives you the right type and sets it up correctly!

- ▶ Possibly wasteful in terms of space
  - ▶ probably no big deal if using text

# Alternative to STI

---

- ▶ STI is good, but what if the different objects have common characteristics but share different representations?
- ▶ Polymorphic association
  - ▶ foreign key that links to objects of different types
  - ▶ along with an indicator of which type
- ▶ articles, sounds, and images all have a catalog\_entry

→ Drops the table first

```
create_table :articles, :force => true do |t|  
  t.column :content, :text  
end
```

```
create_table :sounds, :force => true do |t|  
  t.column :content, :binary  
end
```

```
create_table :images, :force => true do |t|  
  t.column :content, :binary  
end
```

# Polymorphic plumbing

---

```
create_table :catalog_entries, :force => true do |t|
  t.column :name, :string
  t.column :acquired_at, :datetime
  t.column :resource_id, :integer
  t.column :resource_type, :string
end
```

} The combination of these is the foreign key

```
class CatalogEntry < ActiveRecord::Base
  belongs_to :resource, :polymorphic => true
end
```

```
class Article < ActiveRecord::Base
  has_one :catalog_entry, :as => :resource
end
```

ETC.

# Polymorphic examples

---

```
c = CatalogEntry.new(:name => 'Article One' , :acquired_at =>
  Time.now)
c.resource = Article.new(:content => "This is my new article" )
c.save!
```

```
c = CatalogEntry.new(:name => 'Image One' , :acquired_at =>
  Time.now)
c.resource = Image.new(:content => "some binary data" )
c.save!
```

```
c = CatalogEntry.new(:name => 'Sound One' , :acquired_at =>
  Time.now)
c.resource = Sound.new(:content => "more binary data" )
c.save!
```

# What do we get in the DB?

---

```
mysql> select * from articles;
```

id	content
1	This is my new article

```
mysql> select * from images;
```

id	content
1	some binary data

```
mysql> select * from sounds;
```

id	content
1	more binary data

```
mysql> select * from catalog_entries;
```

id	name	acquired_at	resource_id	resource_type
1	Article One	2006-07-18 17:02:05	1	Article
2	Image One	2006-07-18 17:02:05	1	Image
3	Sound One	2006-07-18 17:02:05	1	Sound



# Self referential joins – show demo

---

```
class Employee < ActiveRecord::Base
  belongs_to :manager,
    :class_name => "Employee",
    :foreign_key => "manager_id"

  belongs_to :mentor,
    :class_name => "Employee",
    :foreign_key => "mentor_id"

  has_many :mentored_employees,
    :class_name => "Employee",
    :foreign_key => "mentor_id"

  has_many :managed_employees,
    :class_name => "Employee",
    :foreign_key => "manager_id"
end
```

```
e1 = Employee.new (..)
e2 = Employee.create (..)
e3 = Employee.create (..)
```

```
e1.manager = e2
e1.mentor = e3
e1.save!
```

Queries:

```
e1.manager.name
e1.managed_employees.map {|e|
  e.name}
=> ["dave", "monica"]
```

## Additional structures – *acts as List*

---

Install – gem ‘acts\_as\_list’ in gemfile

```
class TodoList < ActiveRecord::Base
  attr_accessible :name
  has_many :todo_items, :order=>"position"
end
```

```
class TodoItem < ActiveRecord::Base
  attr_accessible :name, :position, :todo_list
  belongs_to :todolist
  acts_as_list :scope => :todo_list
end
```

# Acts\_as\_list

---

```
grocery = TodoList.create(:name=>"grocery")
eggs = TodoItem.create(:name=>"eggs")
eggs #to see the record
milk = TodoItem.create(:name=>"milk")
bread = TodoItem.create(:name=>"bread")

grocery.todo_items << eggs
grocery.todo_items << milk
grocery.todo_items << bread

gli = grocery.todo_items #to see all
gli.first
gli.first.move_to_bottom
grocery.save
```

# Additional structures – *acts as Tree*

---

Sometimes – you need hierarchical relations in the DB

A good example is a category of books

```
create_table :categories, :force => true do |t|  
  t.string :name  
  t.integer :parent_id  
end
```

Table - migration

```
class Category < ActiveRecord::Base  
  acts_as_tree :order => "name"  
end
```

Model

```
root      = Category.create(:name => "Books")  
fiction   = root.children.create(:name => "Fiction")  
non_fiction = root.children.create(:name => "Non Fiction")
```

```
non_fiction.children.create(:name => "Computers")  
non_fiction.children.create(:name => "Science")  
non_fiction.children.create(:name => "Art History")
```

```
fiction.children.create(:name => "Mystery")  
fiction.children.create(:name => "Romance")  
fiction.children.create(:name => "Science Fiction")
```

Populating



# Additional Structures – *Acts as Tree*

```
def display_children(parent)
  puts parent.children(true).map {|child| child.name }.join(", ")
end
```

 `display_children(root)`                      # Fiction, Non Fiction

```
sub_category = root.children.first
puts sub_category.children.size            #=> 3
display_children(sub_category)            #=> Mystery, Romance, Science Fiction
```

```
non_fiction = root.children.find(:first, :conditions => "name = 'Non Fiction'")
```

```
display_children(non_fiction)            #=> Art History, Computers, Science
puts non_fiction.parent.name            #=> Books
```

- ▶ Using the tree.
- ▶ See example and gem at [https://github.com/amerine/acts\\_as\\_tree](https://github.com/amerine/acts_as_tree)

# acts\_as\_commentable (great for posts and comments)

Make your ActiveRecord model act as commentable:

```
class Post < ActiveRecord::Base
  acts_as_commentable
end
```

Add a comment to a model instance:

```
commentable = Post.create
commentable.comments.create(:title => "First comment.", :comment => "This is the first comment.")
```

Fetch comments for a commentable model:

```
commentable = Post.find(1)
comments = commentable.comments.recent.limit(10).all
```

Add multiple type of comments to a model:

```
class Todo < ActiveRecord::Base
  acts_as_commentable :public, :private
end
```

**Note:** This feature is only available from version 4.0 and above

Fetch comments for a this model:

```
public_comments = Todo.find(1).public_comments
private_comments = Todo.find(1).private_comments
```

- [https://github.com/jackdempsey/acts\\_as\\_commentable](https://github.com/jackdempsey/acts_as_commentable)

# Generalizing

---

- ▶ Rails has interesting facilities for ORM
  - ▶ Hibernate, iBATIS, Entity Framework v2, ActiveObjects in Java etc. will be similar in flavor
    - ▶ if not in details
- ▶ Rails is easier to configure
  - ▶ defaults once again eliminate need for configuration files
- ▶ See readings for examples with Hibernate

# Transactions

---

- ▶ **Classic database problem**
  - ▶ transfer \$100 from account 1 to account2
  - ▶ `account1.withdraw(100)`
  - ▶ `account2.deposit(100)`
- ▶ must make sure that both happen or neither happens.
- ▶ Transactions – every SQL statement must succeed or all have no effect.



# ACID

---

- ▶ atomic: all or none
- ▶ consistent: database is in legal state at beginning and at end
- ▶ isolated: no operation outside the transaction can see data in an intermediate state.
  - ▶ transaction history can be serialized.
- ▶ durable – once it's done it can't be undone.

# Transactions at the DB level

---

## ► MySQL

```
start transaction
```

```
update account set balance = balance-100 where  
    account_id=1
```

```
update account set balance = balance+100 where  
account = 2
```

```
commit
```

- some bells and whistles – rollback forces undo
- Transactions are generally database-dependent

# Transactions in rails

---

- ▶ Provided in Active Record

```
Account.transaction (account1, account2) do  
  account1.withdraw(100)  
  account2.deposit(100)  
end
```

- ▶ Rails is smart about dependent/children items – makes them into a transaction

# Race conditions

---

- ▶ Problem with multiple processes accessing a database
  - ▶ as is always the case with web apps
  - ▶ The last one to update *wins* the race
- ▶ p1 sets Order 456 name to “fred”
- ▶ p2 sets Order 456 name to “george”
- ▶ p1 saves
- ▶ p2 saves
- ▶ p1 changes are lost
  
- ▶ Not good.

# Locking

---

- ▶ claim hold on a record until you're done with it
- ▶ Pessimistic locking can be an alternative to transactions
  - ▶ grab hold of everything you need
  - ▶ Do what you need to do
  - ▶ release locks

# Optimistic Locking

---

- ▶ Before you save updated data
  - ▶ make sure row hasn't been changed
- ▶ Rails: Each row has a “version number”
  - ▶ checked before update – if they don't match, abandon & throw exception
- ▶ Enabled if table has a column called `lock_version`.
  - ▶ initialize to zero and leave it alone.

# Optimistic locking example

---

```
create_table :counters, :force => true do |t|  
  t.column :count, :integer  
  t.column :lock_version, :integer, :default => 0  
end
```

```
Counter.create(:count => 0)  
count1 = Counter.find(:first)  
count2 = Counter.find(:first)
```

```
count1.count += 3  
count1.save
```

```
count2.count += 4  
count2.save – this will throw an exception. thinking version 0, but count1.save made version 1
```

- ▶ You have to catch this exception in your code
- ▶ transactions rolled back before exception
- ▶ use this locking and transactions together

# Transaction strategies

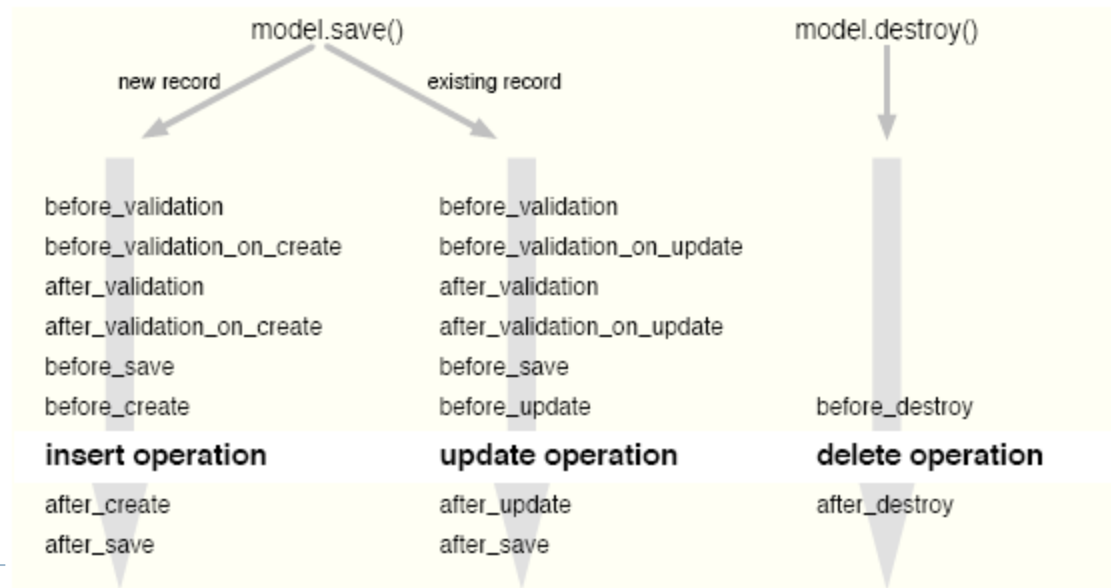
---

- ▶ Note – SQL transactions can include arbitrary tables
- ▶ Rails transactions only one class
  - ▶ Nest them to support more complicated logic
  - ▶ also useful for multi-database systems
- ▶ When to use transactions?
  - ▶ Interacting/dependent objects?
  - ▶ Sequence of changes that makes sense only as a whole



# Callbacks

- ▶ Add code to appropriate points in the life cycle of an object
- ▶ 18 in pairs – before/after
  - ▶ after\_find and after\_initialize



# Example with before\_save

---

- ▶ A simple way to use a callback is to define a method itself.

```
class Order < ActiveRecord::Base
  # ...
  def before_save
    self.payment_due ||= Time.now + 30.days
  end
end
```

# Defining callbacks

---

- ▶ appropriate procedure - `before_save` method for before save event
- ▶ Using handlers
- ▶ or, class to encapsulate callbacks for multiple objects
  - ▶ Say multiple objects need to normalize the credit card number before saving it to the database. Then define a common method and make all models use it.
    - ▶ Another example - encryption

# Observers

---

- ▶ Transparently link to model class
  - ▶ register for callbacks without begin part of model

```
class OrderObserver < ActiveRecord::Observer
  def after_save(an_order)
    an_order.logger.info("Order #{an_order.id} created" )
  end
end
```

- ▶ XXXObserver associated with class XXX
  - ▶ Stored in app\models and call at runtime

- ▶ Install by specifying in application.rb
- ```
config.active_record.observers
  = :order_observer, :audit_observer
```

# Comments on Callbacks and Observers

---

- ▶ Use callbacks to encapsulate common processes in object life cycle
  - ▶ encrypt before save
  - ▶ decrypt after read
- ▶ observer
  - ▶ logging and other subsidiary tasks